

## Creating a Scalable and Highly Available Environment for the Café

### A business request for the café: Implementing a scalable and highly available environment

#### Step 1: Inspecting the environment

In this task, you will evaluate the current state of the environment.

1. Explore the environment, including how the network is set up.

#### Step 2: Creating a NAT gateway for the second Availability Zone

2. Create a NAT gateway in the *Public Subnet* in the second Availability Zone.
3. Configure the network to send internet-bound traffic from instances in *Private Subnet 2* to the NAT gateway you just created.

#### Step 3: Creating a bastion host instance in a public subnet

4. From the **Amazon EC2 console**, create an EC2 instance in one of the public subnets of the *Lab VPC*. It must meet the following criteria:
  - **Name:** **Bastion Host**
  - **Amazon Machine Image (AMI):** *Amazon Linux 2023 AMI*
  - **Instance type:** *t2.micro*
  - Uses the **vockey** key pair
  - **Auto-assign Public IP:** This setting should be enabled
  - Only allows the following traffic:
    - **Type:** *SSH*
    - **Port:** *22*
    - **Source:** Your IP address

#### Step 4: Creating a launch template

5. Create a launch template by using the AMI. It must meet the following criteria.
  - **AMI:** *Cafe WebServer Image*
  - **Instance type:** *t2.micro*

**Key pair (login):** Uses a *new key pair*

- **Security groups:** *CafeSG*
- **Resource tags:**
  - **Key:** *Name*
  - **Value:** *webserver*
  - **Resource types:** *Instances*
- **IAM Instance Profile:** *CafeRole*

#### Step 5: Creating an Auto Scaling group

6. Create a new Auto Scaling Group that meets the following criteria:
  - **Launch template:** Uses the launch template that you created in the previous task

- **VPC:** Uses the VPC that was configured
  - **Subnets:** Uses Private Subnet 1 and Private Subnet 2
  - Skips *all* the advanced options
  - Has a **Group size** configured as:
    - **Desired capacity:** 2
    - **Minimum capacity:** 2
    - **Maximum capacity:** 6
  - Enables the **Target tracking scaling policy** configured as:
    - **Metric type:** *Average CPU utilization*
    - **Target Value:** 25
    - **Instances need:** 60
7. To verify that you created the Auto Scaling group correctly, go to the **Amazon EC2 console**. You should have two instances, both with the name that you configured as *resource tags* in the previous task.

### **Step 6: Creating a load balancer**

8. Create an HTTP Application Load Balancer that meets the following criteria:
- **VPC:** Uses the VPC configured
  - **Subnets:** Uses the two *public subnets*
  - Skips the HTTPS security configuration settings
  - **Security group:** Creates a *new security group* that allows HTTP traffic from anywhere
  - **Target group:** Creates a *new target group*
  - Skips registering targets
9. Modify the Auto Scaling group that you created in the previous task by adding this new load balancer.

### **Step 7: Testing the web application**

10. To test the café web application, visit the Domain Name System (DNS) name of your load balancer and append */cafe* to the URL.

The café application should load.

### **Step 8: Testing automatic scaling under load**

In this task, you will test whether the café application *scales out* automatically.

11. By using *Secure Shell (SSH) passthrough through the bastion host instance*, use SSH to connect to one of the running web server instances.
12. From the web server instance, use the following commands to start a stress test. This test increases the load on the web server CPU:

```
sudo yum install https://dl.fedoraproject.org/pub/epel/epel-release-latest-7.noarch.rpm
sudo yum install stress -y
stress --cpu 1 --timeout 600
```

13. Verify that the Auto Scaling group deploys new instances.
  - Continue to observe the Amazon EC2 console.
  - During the test, you should observe that more web server instances are deployed.